

Introduction to R's time series facilities

Michael Lundholm*

September 22, 2011

Version 1.3

1 Introduction

This is an introduction to R's time series facilities. In an elementary way we deal with reading time series data into R, how time series objects are defined and their properties and finally how time series data can be manipulated. This is *not* an introduction to time series analysis. The note is written as a complement to my colleague Mahmood Arai's "A Brief Guide to R for Beginners in Econometrics" available at http://people.su.se/~ma/R_intro/R_Brief_Guide.pdf. It is presumed that R is installed and that the reader has some basic R-knowledge. An interactive R-session run parallel to reading the notes is recommended.

2 Preliminaries

Time series data have the property of being temporally ordered. Each individual observation have a date and these dates are organised sequentially. We consider here only the case when the time interval between any two observations with sequentially following dates is constant (i.e., the same between all observations). This means that we have a time index $t \in \{1, 2, \dots, T - 1, T\}$, where 1 is the first date and T is the date of the last observation (i.e., the length of the time series). The number of observations per year in a time series is called its *frequency*.

The basic function in R that defines time series is `ts()`. It has as its default frequency `frequency=1`. The starting date is by default year one; `start=1`. The following simple example creates a time series with frequency 1 (yearly data) of the first 10 positive integers, where we have written out the default values:

```
> ts(1:10, start = 1, frequency = 1)
```

*Department of Economics, Stockholm University, michael.lundholm@ne.su.se.

Time Series:

Start = 1

End = 10

Frequency = 1

```
[1] 1 2 3 4 5 6 7 8 9 10
```

Naturally we can set the starting year arbitrarily:

```
> ts(1:10, start = 1998, frequency = 1)
```

Time Series:

Start = 1998

End = 2007

Frequency = 1

```
[1] 1 2 3 4 5 6 7 8 9 10
```

In many instances the interval may be a fraction of a year, say a month or a quarter, so that we 12 or 4 observations per year. That is, the frequency is 12 or 4. Consider quarterly data:

```
> ts(1:10, start = 1998, frequency = 4)
```

	Qtr1	Qtr2	Qtr3	Qtr4
1998	1	2	3	4
1999	5	6	7	8
2000	9	10		

The main point with the function `ts()` is that it defines a time series object which consists of the data and a time line (including frequency). There is no need to define time as a specific variable or use specific time variables in existing data. We will below see how this property can be exploited.

3 Data

For illustrative purposes we will use the Swedish Consumer Price (CPI) with monthly index numbers from January 1980 to December 2005. This data can be downloaded from the Statistics Sweden web page <http://www.scb.se>. However, the data has to be generated from a data base and is therefore also available at <http://people.su.se/~lundh/reproduce/PR0101B1.scb>.¹ Start by downloading the files `PR0101B1.scb` to a working directory at your local computer:²

¹Please note the curious naming convention that the data base of Statistics Sweden is using regarding file names; the base name in upper-case letters and numbers and the extension in lower case letters.

²For details about making data available to R see M. Lundholm, "Loading data into R", Version 1.1. August 18, 2010, http://people.su.se/~lundh/reproduce/loading_data.pdf.

```

> library(utils)
> URL <- "http://people.su.se/~lundh/reproduce/"
> FILE <- "PR0101B1.scb"
> download.file(paste(URL, FILE, sep = ""), FILE)

```

Then load the data file PR0101B1.scb in your favourite text editor and inspect it. The first 10 rows of the file look like this:

```

1 NA
2 1980M01      95.3
3 1980M02      96.8
4 1980M03      97.2
5 1980M04      97.9
6 1980M05      98.2
7 1980M06      98.5
8 1980M07      99.3
9 1980M08      99.9
10 1980M09     102.7

```

NA

```

> cpi <- read.table(FILE, skip = 1, col.names = c("Time",
+         "CPI"))
> head(cpi)

```

```

      Time  CPI
1 1980M01 95.3
2 1980M02 96.8
3 1980M03 97.2
4 1980M04 97.9
5 1980M05 98.2
6 1980M06 98.5

```

```

> str(cpi)

```

```

'data.frame':      312 obs. of  2 variables:
 $ Time: Factor w/ 312 levels "1980M01","1980M02",...: 1 2 3 4 5 6 7 8 9 10 ...
 $ CPI : num  95.3 96.8 97.2 97.9 98.2 ...

```

where the data is assigned the name `cpi`. Note that the resulting object is a data frame.

The command `head(x)` gives the first 6 lines of the object `x` to which it is applied. The dates of the observations are in column 1 which has been assigned the name `Time` and the CPI is in column 2 with the name `CPI`.

4 Creating time series objects with `ts()`

The default of the function `ts()` is `ts(x,start=1,frequency=1)`, where `x` can be a vector, a matrix or a data frame. In order to create a time series object out the object `cpi` (where the first column is redundant) and were the starting date i January 1980 we use the following redefinition:

```
> (cpi1 <- ts(cpi[, 2], start = c(1980, 1), frequency = 12))
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct
1980	95.3	96.8	97.2	97.9	98.2	98.5	99.3	99.9	102.7	104.2
1981	107.2	109.3	109.8	110.5	111.2	111.6	112.6	113.5	114.3	115.0
1982	117.4	119.0	119.3	120.1	120.7	121.1	121.9	122.2	122.9	124.6
1983	129.1	128.8	129.3	130.3	131.1	131.8	132.9	133.5	134.5	135.6
1984	139.4	138.9	140.9	141.8	142.8	142.4	142.8	143.9	144.8	145.5
1985	149.6	151.0	152.1	152.7	154.5	153.9	153.8	153.8	154.5	155.5
1986	158.9	159.0	158.7	159.7	159.7	159.7	160.1	159.9	161.3	161.9
1987	164.4	164.4	164.7	165.1	165.2	164.9	166.9	167.8	169.4	170.1
1988	171.6	172.9	173.7	175.2	175.8	176.3	177.1	177.5	178.8	180.2
1989	183.0	184.0	184.7	186.5	187.3	187.9	187.9	188.7	190.2	191.8
1990	199.0	199.9	205.4	205.2	206.4	206.2	208.2	209.6	212.0	213.4
1991	218.9	225.0	225.8	227.1	227.3	227.0	227.1	226.7	229.2	230.1
1992	230.2	230.3	231.3	231.9	232.0	231.5	231.2	231.3	234.6	235.1
1993	241.0	241.6	242.7	243.7	243.1	242.3	241.9	242.3	244.5	245.2
1994	245.1	245.9	246.8	247.8	248.3	248.4	248.4	248.5	250.7	251.0
1995	251.3	252.3	253.3	255.0	255.3	255.1	254.8	254.5	256.2	256.9
1996	255.6	255.8	257.0	257.6	257.3	256.3	255.7	254.5	256.0	255.9
1997	254.6	254.2	255.2	257.0	257.0	257.4	257.3	257.4	259.8	259.6
1998	256.9	256.6	257.0	257.7	258.1	257.6	257.0	255.7	256.8	257.3
1999	256.2	256.3	257.3	257.9	258.3	258.7	257.6	257.6	259.4	259.7
2000	257.5	258.7	259.9	260.0	261.3	261.2	260.0	260.2	262.0	262.6
2001	261.7	262.6	264.6	266.9	268.7	268.3	266.9	267.6	269.9	269.1
2002	268.8	269.4	271.8	272.9	273.6	273.2	272.3	272.4	274.5	275.4
2003	276.0	278.4	279.8	278.8	278.5	277.7	276.8	276.7	278.7	278.9
2004	278.0	277.3	279.4	279.4	280.1	278.9	278.5	278.2	280.2	281.0
2005	277.9	279.2	279.8	280.2	280.3	280.4	279.4	279.9	281.9	282.4
	Nov	Dec								
1980	104.8	105.2								
1981	115.4	114.9								
1982	125.6	125.9								
1983	136.4	137.5								
1984	146.4	148.8								
1985	156.5	157.1								
1986	161.9	162.3								
1987	170.7	170.7								

```

1988 180.5 180.9
1989 192.2 192.8
1990 214.1 213.9
1991 231.1 230.8
1992 234.0 234.9
1993 245.3 244.3
1994 250.8 250.4
1995 256.8 256.0
1996 255.3 254.9
1997 259.2 259.1
1998 256.7 256.2
1999 259.0 259.6
2000 262.7 262.5
2001 269.2 269.5
2002 274.7 275.1
2003 278.3 278.6
2004 279.4 279.4
2005 281.7 281.8

```

```
> str(cpi1)
```

```
Time-Series [1:312] from 1980 to 2006: 95.3 96.8 97.2 97.9 98.2 ...
```

Note first that `cpi[,2]` picks the entire second column. The argument `start=c(1980,1)` specifies the starting date as the date of the first observation in 1980. Since `frequency=12` defines data as monthly, the first observation in 1980 is January 1980. The resulting object is a time series object.

The `ts()` function has two clones: `is.ts()` which checks whether an object is a time series object and `as.ts()` which coerces an object to be a time series object:

```
> is.ts(cpi1)
```

```
[1] TRUE
```

```
> is.ts(cpi)
```

```
[1] FALSE
```

```
> is.ts(as.ts(cpi))
```

```
[1] TRUE
```

Note `cpi` is not a time series object to R even if it contains information about the dates of the observation. To R these dates are just factors.

Note finally that even if `cpi1` is a time series object, a part of this object selected through indexing loses these properties:

```
> cpi1[7:18]

[1] 99.3 99.9 102.7 104.2 104.8 105.2 107.2 109.3 109.8 110.5
[11] 111.2 111.6

> is.ts(cpi1[7:18])

[1] FALSE
```

i.e., the latter being an attempt to access data from July 1980 to June 1981. Well, we do but the resulting object does not inherit the relevant time series properties. We will have to use other methods to consider subsets of time series objects.

5 Extracting attributes from time series objects

There are several commands that can extract different attributes from a time series object. Of course one can extract the time line (i.e., the dates of the observations):

```
> cpi1.time <- time(cpi1)
> head(cpi1.time)

[1] 1980.0 1980.1 1980.2 1980.2 1980.3 1980.4
```

We see that the time index is given as decimal numbers. The difference between two points in time (time difference) is a fixed number which depends on the frequency (i.e., it is the inverse of the frequency). The time difference between two observations in a time series object and the frequency of the object can of course be extracted:

```
> deltat(cpi1)

[1] 0.083333

> frequency(cpi1)

[1] 12
```

and we verify that $\frac{1}{12} \approx 0.083$. Related to the frequency is of course where in a cycle (i.e., in which position during a period/year a certain observations have). Consider the following example:

```
> (xseries <- ts(matrix(rep(2, 12)), start = c(1980,
+      3), frequency = 12))
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1980			2	2	2	2	2	2	2	2	2	2
1981	2	2										

```
> cycle(xseries)
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1980			3	4	5	6	7	8	9	10	11	12
1981	1	2										

That is, the time series `xseries` consists only of twelve 2's with starting date March 1980. It is monthly data since frequency is 12. The first observation has number 3 in the cycle (March being the third month of the year). The last observation (the 12'th) is second in the cycle since it has date February 1981.

Finally, we can extract the start and end dates of a time series object:

```
> start(cpi1)
```

```
[1] 1980    1
```

```
> end(cpi1)
```

```
[1] 2005   12
```

6 Selecting subsets of time series objects

We previously noted that although `cpi1` may be a time series object, say, `cpi1[221]` is not. So how can we define subsets of time series objects?³ This is done with the function `window()`. The function takes `start` and `end` as arguments. Say we want to extract the period July 1987 to June 1989 from `cpi1`:

```
> window(cpi1, start = c(1987, 7), end = c(1988, 6))
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct
1987							166.9	167.8	169.4	170.1
1988	171.6	172.9	173.7	175.2	175.8	176.3				
	Nov	Dec								
1987	170.7	170.7								
1988										

If the argument `frequency` and/or `deltat` is provided, then the resulting time series is re-sampled at the new frequency given; else frequency is inherited from the original time series object. Hence,

³`subset()` cannot be used since it does not take time series objects as arguments.

```
> window(cpi1, start = c(1980, 2), deltat = 1)
```

Time Series:

Start = 1980.1

End = 2005.1

Frequency = 1

```
[1] 96.8 109.3 119.0 128.8 138.9 151.0 159.0 164.4 172.9 184.0
[11] 199.9 225.0 230.3 241.6 245.9 252.3 255.8 254.2 256.6 256.3
[21] 258.7 262.6 269.4 278.4 277.3 279.2
```

re-samples data with frequency 1 starting with February 1980; i.e., the data from February each year will be extracted to constitute a new time series but now re-samples as yearly data (one observation per period/year).

We split the data into two time series objects:

```
> cpi1.a <- window(cpi1, end = c(1993, 12))
```

```
> cpi1.b <- window(cpi1, start = c(1992, 1))
```

The time series objects `cpi1.a` and `cpi1.b` now constitutes two different parts of the original time series object `cpi1`. They are, however, not mutually exclusive since data for the years 1992–1993 are in both. We can now combine these two objects using `ts.intersect()` and `ts.union()`:

```
> cpi1.c <- ts.union(cpi1.a, cpi1.b)
```

```
> cpi1.d <- ts.intersect(cpi1.a, cpi1.b)
```

We make the following observations regarding the created objects `cpi1.c` and `cpi1.d`:

- They are both time series objects containing two variables (`cpi1.a` and `cpi1.b`).
- `cpi1.c` is the union between `cpi1.a` and `cpi1.b` and contains 312 observations of the two variables. The first date is January 1980 and the last December 2005. There are a lot of missing values (NA):

```
> str(cpi1.c)
```

```
mts [1:312, 1:2] 95.3 96.8 97.2 97.9 98.2 ...
- attr(*, "dimnames")=List of 2
..$ : NULL
..$ : chr [1:2] "cpi1.a" "cpi1.b"
- attr(*, "tsp")= num [1:3] 1980 2006 12
- attr(*, "class")= chr [1:2] "mts" "ts"
```

```
> summary(cpi1.c)
```


cpi1.a	cpi1.b
Min. : 95.3	Min. :230
1st Qu.:132.6	1st Qu.:255
Median :163.3	Median :258
Mean :169.0	Mean :260
3rd Qu.:206.8	3rd Qu.:273
Max. :245.3	Max. :282
NA's :144.0	NA's :144

- `cpi1.d` is the intersection between `cpi1.a` and `cpi1.b` and contains only 24 observations of the two variables and where the variables are identical. The first date is January 1992 and last date December 1993. There are no missing values.

```
> str(cpi1.d)

mts [1:24, 1:2] 230 230 231 232 232 ...
- attr(*, "dimnames")=List of 2
..$ : NULL
..$ : chr [1:2] "cpi1.a" "cpi1.b"
- attr(*, "tsp")= num [1:3] 1992 1994 12
- attr(*, "class")= chr [1:2] "mts" "ts"

> summary(cpi1.d)
```

cpi1.a	cpi1.b
Min. :230	Min. :230
1st Qu.:232	1st Qu.:232
Median :238	Median :238
Mean :238	Mean :238
3rd Qu.:243	3rd Qu.:243
Max. :245	Max. :245

7 Lags and differences

In time series analysis observations from different dates, say y_t from date t and the preceding observation y_{t-1} from date $t-1$, are often used in the same model. In terms of terminology and mathematical notation y_{t-1} one usually call the (one period) lag of y_t . Sometimes a lag operator L is introduced, such that $L y_t = y_{t-1}$. For longer lags, say lag n , one write $L^n y_t = y_{t-n}$.

In R lags are created by the function `lag(y,k=1)`, where the default `k=1` implies a *lead*. Therefore, note that the default value is equivalent to $L^{-1} y_t = y_{t+1}$! To get $L y_t = y_{t-1}$ we must write `lag(y,k=-1)`. Consider the following example which implements y_t and $L y_t = y_{t-1}$:

```
> ts(1:5)
```

```
Time Series:
```

```
Start = 1
```

```
End = 5
```

```
Frequency = 1
```

```
[1] 1 2 3 4 5
```

```
> lag(ts(1:5), k = -1)
```

```
Time Series:
```

```
Start = 2
```

```
End = 6
```

```
Frequency = 1
```

```
[1] 1 2 3 4 5
```

In `ts(1:5)` the fifth observation with value 5 has date 5. On the other hand, in `lag(ts(1:5),k=-1)` the fifth observation with value 5 has date 6. Therefore, only the time index is changed! The vector of observations is unaffected, which we see if we pick (say) the third observation of the time series and its lag:

```
> ts(1:5)[3]
```

```
[1] 3
```

```
> lag(ts(1:5), k = -1)[3]
```

```
[1] 3
```

This means that one has to be careful using a time series and its lag in the same regression; see below.

R also has a function which directly defines differences between variables of different dates. This is the function `diff(y,lag=1,difference=1)`, with default arguments. Here the option `lag` has its intuitive meaning. The default values give the equivalent to `y-lag(y,k=-1)`:

```
> ts(1:5) - lag(ts(1:5), k = -1)
```

```
Time Series:
```

```
Start = 2
```

```
End = 5
```

```
Frequency = 1
```

```
[1] 1 1 1 1
```

```
> diff(ts(1:5), lag = 1, difference = 1)
```

```

Time Series:
Start = 2
End = 5
Frequency = 1
[1] 1 1 1 1

```

This means that `diff(y,lag=1,difference=1)` is exactly equivalent to the general difference operator Δ , which is defined as $\Delta y_t = (1-L)y_t = y_t - y_{t-1}$.

However, one need also to be careful using `diff()`. For instance $y_t - y_{t-1}$ is implemented by

```
> diff(ts(1:5), lag = 2, difference = 1)
```

```

Time Series:
Start = 3
End = 5
Frequency = 1
[1] 2 2 2

```

which is replicate using `lag` as

```
> ts(1:5) - lag(ts(1:5), k = -2)
```

```

Time Series:
Start = 3
End = 5
Frequency = 1
[1] 2 2 2

```

On the other hand

```
> diff(ts(1:5), lag = 1, difference = 2)
```

```

Time Series:
Start = 3
End = 5
Frequency = 1
[1] 0 0 0

```

implements $\Delta^2 y_t = (1-L^2)y_t = y_t - 2y_{t-1} + y_{t-2}$ which in turn is replicated by

```

> ts(1:5) - 2 * lag(ts(1:5), k = -1) + lag(ts(1:5),
+      k = -2)

```

```

Time Series:
Start = 3
End = 5
Frequency = 1
[1] 0 0 0

```

Finally, the function `diffinv()`. Given some vector of difference and an intimal value the undifferentiated series can be retrieved. Consider the first order difference $\Delta y_t = y_t - y_{t-1}$. Given y_1 and Δy_t we can calculate y_t using $y_t = y_{t-1} + \Delta y_t$ or

```

> diffinv(diff(ts(1:5), lag = 1, difference = 1), lag = 1,
+         difference = 1, 1)

```

```

Time Series:
Start = 1
End = 5
Frequency = 1
[1] 1 2 3 4 5

```

where the last 1 is y_1 (the first observation in the original series).

8 Regression and time series data

Let us define $y_t = \Delta \text{CPI}_t$ and suppose that we want estimate the linear model $y_t = \alpha_0 + \alpha_1 y_{t-1}$. Note that the regression coefficient α_1 in such model is identical to the coefficient of correlation between y_t and y_{t-1} . If we plot `diff(cpi1)` against time we get an idea about the correlation:

```

> plot(diff(cpi1))

```

The correlation between the two series is accordingly very small. However, calculating the coefficient of correlation explicitly we get

```
> cor(diff(cpi1), lag(diff(cpi1), k = -1))
```

```
[1] 1
```

which is a result that we also get if we estimate the regression model with `lm()`:

```
> result <- lm(diff(cpi1) ~ lag(diff(cpi1), -1))
> summary(result)
```

Call:

```
lm(formula = diff(cpi1) ~ lag(diff(cpi1), -1))
```

Residuals:

	Min	1Q	Median	3Q	Max
	-7.36e-15	1.00e-18	1.80e-17	3.70e-17	3.39e-16

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-3.02e-16	2.73e-17	-1.11e+01	<2e-16 ***
lag(diff(cpi1), -1)	1.00e+00	2.14e-17	4.67e+16	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
Residual standard error: 4.25e-16 on 309 degrees of freedom
Multiple R-squared: 1, Adjusted R-squared: 1
F-statistic: 2.18e+33 on 1 and 309 DF, p-value: <2e-16
```

Despite that we can see in the graph that the correlation is low we get in both these instances that the coefficient of correlation is exactly unity.

The explanation is that `cor()` and `lm()` (as well as many other functions) are completely unaware of the time series properties of the data set. For instance, `lm()` takes as data argument a data frame or something that can be coerced to be a data frame. When our data is coerced to become a data frame it loses its time series properties:

```
> z <- data.frame(diff(cpi1), lag(diff(cpi1), -1))
> is.ts(z)
```

```
[1] FALSE
```

This can also be seen by looking at the first elements of `diff(cpi1)` and its lag by use of indices:

```
> diff(cpi1)[1:10]
```

```
[1] 1.5 0.4 0.7 0.3 0.3 0.8 0.6 2.8 1.5 0.6
```

```
> lag(diff(cpi1), -1)[1:10]
```

```
[1] 1.5 0.4 0.7 0.3 0.3 0.8 0.6 2.8 1.5 0.6
```

They are identical even though the observations refer to different dates. This means that the (standard) method of just imputing the mathematical expressions into the formula to be estimated does not work with the lag and difference operators.

Now, there are at least two methods to overcome this. The first method is to create a data set such that the variables get a common time index:

```
> library(tseries)
> y <- ts.union(diff(cpi1), lag(diff(cpi1), -1))
> y <- na.remove(y)
> is.ts(y[, 1])
```

```
[1] TRUE
```

```
> is.ts(y[, 2])
```

```
[1] TRUE
```

Note the now the separate columns of the time series object `y` are time series objects in themselves. This means that rows in the object have a common date. The library `tsseries` is loaded to get access to the function `na.remove()`, which (as indicated by the name) removes missing values from time series objects. This accounts for the lower degrees of freedom below. Now the model is estimated

```
> result <- lm(y[, 1] ~ y[, 2])
> summary(result)
```

Call:

```
lm(formula = y[, 1] ~ y[, 2])
```

Residuals:

Min	1Q	Median	3Q	Max
-2.697	-0.707	-0.096	0.463	5.603

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.4947	0.0717	6.90	2.9e-11 ***
y[, 2]	0.1698	0.0561	3.03	0.0027 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.11 on 308 degrees of freedom

Multiple R-squared: 0.0289, Adjusted R-squared: 0.0257

F-statistic: 9.16 on 1 and 308 DF, p-value: 0.00269

Now it works because even if data with the common time index is coerced to a data frame its two vectors are not identical:

```
> y[1:10, 1]
```

```
[1] 0.4 0.7 0.3 0.3 0.8 0.6 2.8 1.5 0.6 0.4
```

```
> y[1:10, 2]
```

```
[1] 1.5 0.4 0.7 0.3 0.3 0.8 0.6 2.8 1.5 0.6
```

We can also use `cor()` to get the same result:

```
> cor(y[, 1], y[, 2])
```

```
[1] 0.1699
```

The second method is to invoke the library `dyn`. From the manual of the library we read:

“**dyn** enables regression functions that were not written to handle time series to handle them. Both the dependent and independent variables may be time series and they may have different time indexes (in which case they are automatically aligned). The time series may also have missing values including internal missing values.”

```
> library(dyn)
> result <- dyn$lm(diff(cpi1) ~ lag(diff(cpi1), -1))
> summary(result)
```

Call:
lm(formula = dyn(diff(cpi1) ~ lag(diff(cpi1), -1)))

Residuals:

Min	1Q	Median	3Q	Max
-2.697	-0.707	-0.096	0.463	5.603

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.4947	0.0717	6.90	2.9e-11 ***
lag(diff(cpi1), -1)	0.1698	0.0561	3.03	0.0027 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.11 on 308 degrees of freedom
(2 observations deleted due to missingness)
Multiple R-squared: 0.0289, Adjusted R-squared: 0.0257
F-statistic: 9.16 on 1 and 308 DF, p-value: 0.00269

We verify that the two methods give the same result; $\hat{\alpha}_1 = 0.17$. One difference, though, is that although in both cases we have 308 degrees of freedom we get using **dyn** a report of missing values (2 missing values in fact). The reason is that this is in relation to original (undifferentiated and unlagged) the time series object `cpi1`, which we have differenced (one missing value) and lagged (another missing value).

9 Time and seasonal dummy variables

Once we have defined a time series using `ts()` it is easy to define time and seasonal dummy variables. We start with the time dummy using `time()`:

```
> time(cpi1)
```


	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug
1980	1980.0	1980.1	1980.2	1980.2	1980.3	1980.4	1980.5	1980.6
1981	1981.0	1981.1	1981.2	1981.2	1981.3	1981.4	1981.5	1981.6
1982	1982.0	1982.1	1982.2	1982.2	1982.3	1982.4	1982.5	1982.6
1983	1983.0	1983.1	1983.2	1983.2	1983.3	1983.4	1983.5	1983.6
1984	1984.0	1984.1	1984.2	1984.2	1984.3	1984.4	1984.5	1984.6
1985	1985.0	1985.1	1985.2	1985.2	1985.3	1985.4	1985.5	1985.6
1986	1986.0	1986.1	1986.2	1986.2	1986.3	1986.4	1986.5	1986.6
1987	1987.0	1987.1	1987.2	1987.2	1987.3	1987.4	1987.5	1987.6
1988	1988.0	1988.1	1988.2	1988.2	1988.3	1988.4	1988.5	1988.6
1989	1989.0	1989.1	1989.2	1989.2	1989.3	1989.4	1989.5	1989.6
1990	1990.0	1990.1	1990.2	1990.2	1990.3	1990.4	1990.5	1990.6
1991	1991.0	1991.1	1991.2	1991.2	1991.3	1991.4	1991.5	1991.6
1992	1992.0	1992.1	1992.2	1992.2	1992.3	1992.4	1992.5	1992.6
1993	1993.0	1993.1	1993.2	1993.2	1993.3	1993.4	1993.5	1993.6
1994	1994.0	1994.1	1994.2	1994.2	1994.3	1994.4	1994.5	1994.6
1995	1995.0	1995.1	1995.2	1995.2	1995.3	1995.4	1995.5	1995.6
1996	1996.0	1996.1	1996.2	1996.2	1996.3	1996.4	1996.5	1996.6
1997	1997.0	1997.1	1997.2	1997.2	1997.3	1997.4	1997.5	1997.6
1998	1998.0	1998.1	1998.2	1998.2	1998.3	1998.4	1998.5	1998.6
1999	1999.0	1999.1	1999.2	1999.2	1999.3	1999.4	1999.5	1999.6
2000	2000.0	2000.1	2000.2	2000.2	2000.3	2000.4	2000.5	2000.6
2001	2001.0	2001.1	2001.2	2001.2	2001.3	2001.4	2001.5	2001.6
2002	2002.0	2002.1	2002.2	2002.2	2002.3	2002.4	2002.5	2002.6
2003	2003.0	2003.1	2003.2	2003.2	2003.3	2003.4	2003.5	2003.6
2004	2004.0	2004.1	2004.2	2004.2	2004.3	2004.4	2004.5	2004.6
2005	2005.0	2005.1	2005.2	2005.2	2005.3	2005.4	2005.5	2005.6
	Sep	Oct	Nov	Dec				
1980	1980.7	1980.8	1980.8	1980.9				
1981	1981.7	1981.8	1981.8	1981.9				
1982	1982.7	1982.8	1982.8	1982.9				
1983	1983.7	1983.8	1983.8	1983.9				
1984	1984.7	1984.8	1984.8	1984.9				
1985	1985.7	1985.8	1985.8	1985.9				
1986	1986.7	1986.8	1986.8	1986.9				
1987	1987.7	1987.8	1987.8	1987.9				
1988	1988.7	1988.8	1988.8	1988.9				
1989	1989.7	1989.8	1989.8	1989.9				
1990	1990.7	1990.8	1990.8	1990.9				
1991	1991.7	1991.8	1991.8	1991.9				
1992	1992.7	1992.8	1992.8	1992.9				
1993	1993.7	1993.8	1993.8	1993.9				
1994	1994.7	1994.8	1994.8	1994.9				
1995	1995.7	1995.8	1995.8	1995.9				

```

1996 1996.7 1996.8 1996.8 1996.9
1997 1997.7 1997.8 1997.8 1997.9
1998 1998.7 1998.8 1998.8 1998.9
1999 1999.7 1999.8 1999.8 1999.9
2000 2000.7 2000.8 2000.8 2000.9
2001 2001.7 2001.8 2001.8 2001.9
2002 2002.7 2002.8 2002.8 2002.9
2003 2003.7 2003.8 2003.8 2003.9
2004 2004.7 2004.8 2004.8 2004.9
2005 2005.7 2005.8 2005.8 2005.9

```

Note that (i) the variable is a time series variable and (ii) that units are in increments of a year starting at 1980. However, since we most frequently use time variables to remove time trends units may matter. If units matter we can transform units to something appropriate. For example

```
> (TIME <- 1 + frequency(cpi1) * (time(cpi1) - start(cpi1)[1]))
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1980	1	2	3	4	5	6	7	8	9	10	11	12
1981	13	14	15	16	17	18	19	20	21	22	23	24
1982	25	26	27	28	29	30	31	32	33	34	35	36
1983	37	38	39	40	41	42	43	44	45	46	47	48
1984	49	50	51	52	53	54	55	56	57	58	59	60
1985	61	62	63	64	65	66	67	68	69	70	71	72
1986	73	74	75	76	77	78	79	80	81	82	83	84
1987	85	86	87	88	89	90	91	92	93	94	95	96
1988	97	98	99	100	101	102	103	104	105	106	107	108
1989	109	110	111	112	113	114	115	116	117	118	119	120
1990	121	122	123	124	125	126	127	128	129	130	131	132
1991	133	134	135	136	137	138	139	140	141	142	143	144
1992	145	146	147	148	149	150	151	152	153	154	155	156
1993	157	158	159	160	161	162	163	164	165	166	167	168
1994	169	170	171	172	173	174	175	176	177	178	179	180
1995	181	182	183	184	185	186	187	188	189	190	191	192
1996	193	194	195	196	197	198	199	200	201	202	203	204
1997	205	206	207	208	209	210	211	212	213	214	215	216
1998	217	218	219	220	221	222	223	224	225	226	227	228
1999	229	230	231	232	233	234	235	236	237	238	239	240
2000	241	242	243	244	245	246	247	248	249	250	251	252
2001	253	254	255	256	257	258	259	260	261	262	263	264
2002	265	266	267	268	269	270	271	272	273	274	275	276
2003	277	278	279	280	281	282	283	284	285	286	287	288
2004	289	290	291	292	293	294	295	296	297	298	299	300
2005	301	302	303	304	305	306	307	308	309	310	311	312

which also preserves the time series properties. Note that both these examples are independent of the length of the time series variable and its frequency, which can be useful in programming. Equivalent to the second example in all respects is the following:

```
> (TIME2 <- ts(1:length(cpi1), start = start(cpi1),
+ frequency = frequency(cpi1)))
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1980	1	2	3	4	5	6	7	8	9	10	11	12
1981	13	14	15	16	17	18	19	20	21	22	23	24
1982	25	26	27	28	29	30	31	32	33	34	35	36
1983	37	38	39	40	41	42	43	44	45	46	47	48
1984	49	50	51	52	53	54	55	56	57	58	59	60
1985	61	62	63	64	65	66	67	68	69	70	71	72
1986	73	74	75	76	77	78	79	80	81	82	83	84
1987	85	86	87	88	89	90	91	92	93	94	95	96
1988	97	98	99	100	101	102	103	104	105	106	107	108
1989	109	110	111	112	113	114	115	116	117	118	119	120
1990	121	122	123	124	125	126	127	128	129	130	131	132
1991	133	134	135	136	137	138	139	140	141	142	143	144
1992	145	146	147	148	149	150	151	152	153	154	155	156
1993	157	158	159	160	161	162	163	164	165	166	167	168
1994	169	170	171	172	173	174	175	176	177	178	179	180
1995	181	182	183	184	185	186	187	188	189	190	191	192
1996	193	194	195	196	197	198	199	200	201	202	203	204
1997	205	206	207	208	209	210	211	212	213	214	215	216
1998	217	218	219	220	221	222	223	224	225	226	227	228
1999	229	230	231	232	233	234	235	236	237	238	239	240
2000	241	242	243	244	245	246	247	248	249	250	251	252
2001	253	254	255	256	257	258	259	260	261	262	263	264
2002	265	266	267	268	269	270	271	272	273	274	275	276
2003	277	278	279	280	281	282	283	284	285	286	287	288
2004	289	290	291	292	293	294	295	296	297	298	299	300
2005	301	302	303	304	305	306	307	308	309	310	311	312

All three can be used in regressions using both methods from the previous section:

```
> summary(dyn$lm(diff(cpi1) ~ lag(diff(cpi1), -1) +
+ time(cpi1)))
```

Call:

```
lm(formula = dyn(diff(cpi1) ~ lag(diff(cpi1), -1) + time(cpi1)))
```

```

Residuals:
    Min       1Q   Median       3Q      Max
-2.556 -0.700 -0.177  0.433  5.507

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    63.78351   17.06438    3.74 0.00022 ***
lag(diff(cpi1), -1) 0.11918    0.05665    2.10 0.03621 *
time(cpi1)     -0.03174    0.00856   -3.71 0.00025 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.09 on 307 degrees of freedom
(3 observations deleted due to missingness)
Multiple R-squared: 0.0705,    Adjusted R-squared: 0.0645
F-statistic: 11.6 on 2 and 307 DF,  p-value: 1.33e-05

> summary(dyn$lm(diff(cpi1) ~ lag(diff(cpi1), -1) +
+      TIME))

Call:
lm(formula = dyn(diff(cpi1) ~ lag(diff(cpi1), -1) + TIME))

Residuals:
    Min       1Q   Median       3Q      Max
-2.556 -0.700 -0.177  0.433  5.507

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    0.941693   0.139484    6.75 7.3e-11 ***
lag(diff(cpi1), -1) 0.119183   0.056652    2.10 0.03621 *
TIME           -0.002645   0.000713   -3.71 0.00025 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.09 on 307 degrees of freedom
(3 observations deleted due to missingness)
Multiple R-squared: 0.0705,    Adjusted R-squared: 0.0645
F-statistic: 11.6 on 2 and 307 DF,  p-value: 1.33e-05

> summary(dyn$lm(diff(cpi1) ~ lag(diff(cpi1), -1) +
+      TIME2))

Call:
lm(formula = dyn(diff(cpi1) ~ lag(diff(cpi1), -1) + TIME2))

```

Residuals:

Min	1Q	Median	3Q	Max
-2.556	-0.700	-0.177	0.433	5.507

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.941693	0.139484	6.75	7.3e-11 ***
lag(diff(cpi1), -1)	0.119183	0.056652	2.10	0.03621 *
TIME2	-0.002645	0.000713	-3.71	0.00025 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.09 on 307 degrees of freedom

(3 observations deleted due to missingness)

Multiple R-squared: 0.0705, Adjusted R-squared: 0.0645

F-statistic: 11.6 on 2 and 307 DF, p-value: 1.33e-05

In terms of inference they are equivalent, but (of course) the choice time units affect the size of regression coefficients and their estimated standard deviations.

Seasonal dummy variables are more complicated but nonetheless easier to construct. Basically, the number of seasons equals the frequency in the data and we therefore need **frequency** minus one seasonal dummy variables. We show only one method from package **forecast** which is really simple:

```
> library(forecast)
```

This is forecast 3.03

```
> SEASON <- ts(seasonaldummy(diff(cpi1)), start = start(diff(cpi1)),  
+ frequency = frequency(diff(cpi1)))  
> summary(dyn$lm(diff(cpi1) ~ SEASON))
```

Call:

```
lm(formula = dyn(diff(cpi1) ~ SEASON))
```

Residuals:

Min	1Q	Median	3Q	Max
-3.320	-0.485	-0.069	0.431	5.215

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.1692	0.1960	0.86	0.38850
SEASON1	0.9508	0.2799	3.40	0.00077 ***

SEASON2	0.7154	0.2771	2.58	0.01032 *
SEASON3	0.9808	0.2771	3.54	0.00047 ***
SEASON4	0.6538	0.2771	2.36	0.01895 *
SEASON5	0.3385	0.2771	1.22	0.22293
SEASON6	-0.3154	0.2771	-1.14	0.25601
SEASON7	-0.1654	0.2771	-0.60	0.55110
SEASON8	0.0385	0.2771	0.14	0.88971
SEASON9	1.6000	0.2771	5.77	1.9e-08 ***
SEASON10	0.5308	0.2771	1.92	0.05641 .
SEASON11	-0.1423	0.2771	-0.51	0.60797

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.999 on 299 degrees of freedom

Multiple R-squared: 0.242, Adjusted R-squared: 0.214

F-statistic: 8.66 on 11 and 299 DF, p-value: 2.56e-13

Note that with monthly data only 11 variables are created. Note also that the number of observations is 311 due to the first order difference and for this reason the definition of the dummy variable is based on the first order difference.

10 Aggregation of time series data

Sometimes data are available as high frequency data but what you need is low frequency data. Say, you have monthly data but need quarterly or yearly data. The function `aggregate()` has a method for time series objects so that one can use `aggregate(x, nfrequency = 1, FUN = sum, ...)`, where `x` is a time series object. The function splits `x` in blocks of length `frequency(x)/nfrequency`, where `nfrequency` is the frequency of the aggregated data (which must be lower than the frequency of `x` and which is 1 (yearly data) by default). The function specified by the argument `FUN` is then applied to each block of data (be default `sum`). For other arguments see `?aggregate`.

The default function `sum` may be suitable for flow data as (for example) the data set `AirPassengers`, which records the total number of international air passengers per month. Say we want to aggregate this data to quarterly data

```
> AirPassengers
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1949	112	118	132	129	121	135	148	148	136	119	104	118
1950	115	126	141	135	125	149	170	170	158	133	114	140

```

1951 145 150 178 163 172 178 199 199 184 162 146 166
1952 171 180 193 181 183 218 230 242 209 191 172 194
1953 196 196 236 235 229 243 264 272 237 211 180 201
1954 204 188 235 227 234 264 302 293 259 229 203 229
1955 242 233 267 269 270 315 364 347 312 274 237 278
1956 284 277 317 313 318 374 413 405 355 306 271 306
1957 315 301 356 348 355 422 465 467 404 347 305 336
1958 340 318 362 348 363 435 491 505 404 359 310 337
1959 360 342 406 396 420 472 548 559 463 407 362 405
1960 417 391 419 461 472 535 622 606 508 461 390 432

```

```
> (AirPassengers.Q <- aggregate(AirPassengers, nfrequency = 4))
```

```

      Qtr1 Qtr2 Qtr3 Qtr4
1949   362  385  432  341
1950   382  409  498  387
1951   473  513  582  474
1952   544  582  681  557
1953   628  707  773  592
1954   627  725  854  661
1955   742  854 1023  789
1956   878 1005 1173  883
1957   972 1125 1336  988
1958  1020 1146 1400 1006
1959  1108 1288 1570 1174
1960  1227 1468 1736 1283

```

That is, the monthly observations are just summed quarter by quarter.

Note however, that you need monthly observations for “full” quarters for this result. Say data is incomplete in the sense that it starts in February:

```

> AirPassengers.N <- window(AirPassengers, start = c(1949,
+ 2))
> (AirPassengers.N.Q <- aggregate(AirPassengers.N,
+ nfrequency = 4))

```

Time Series:

Start = 1949.083333333333

End = 1960.583333333333

Frequency = 4

```

[1] 379 404 403 337 402 444 461 399 491 549 545 483
[13] 554 631 642 562 667 736 720 585 650 800 781 674
[25] 769 949 933 799 907 1105 1066 892 1005 1242 1218 981
[37] 1028 1289 1268 1007 1144 1440 1429 1184 1271 1629 1575

```

Aggregation will lead to quarters, but quarters consisting of February–April, May–July, August–October and November–January. The last two observations in the data set `AirPassengers.N`, November and December 1960, do not constitute a full quarter and are therefore dismissed.

If data are flows the default function `sum` is appropriate, but one may also have data in levels. Then (say) `mean` is an alternative in aggregation. The data set `nottem` contains monthly average air temperatures at Nottingham Castle in degrees Fahrenheit for the period 1920–1939. Say we want to construct quarterly averages:

```
> nottem
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1920	40.6	40.8	44.4	46.7	54.1	58.5	57.7	56.4	54.3	50.5	42.9	39.8
1921	44.2	39.8	45.1	47.0	54.1	58.7	66.3	59.9	57.0	54.2	39.7	42.8
1922	37.5	38.7	39.5	42.1	55.7	57.8	56.8	54.3	54.3	47.1	41.8	41.7
1923	41.8	40.1	42.9	45.8	49.2	52.7	64.2	59.6	54.4	49.2	36.3	37.6
1924	39.3	37.5	38.3	45.5	53.2	57.7	60.8	58.2	56.4	49.8	44.4	43.6
1925	40.0	40.5	40.8	45.1	53.8	59.4	63.5	61.0	53.0	50.0	38.1	36.3
1926	39.2	43.4	43.4	48.9	50.6	56.8	62.5	62.0	57.5	46.7	41.6	39.8
1927	39.4	38.5	45.3	47.1	51.7	55.0	60.4	60.5	54.7	50.3	42.3	35.2
1928	40.8	41.1	42.8	47.3	50.9	56.4	62.2	60.5	55.4	50.2	43.0	37.3
1929	34.8	31.3	41.0	43.9	53.1	56.9	62.5	60.3	59.8	49.2	42.9	41.9
1930	41.6	37.1	41.2	46.9	51.2	60.4	60.1	61.6	57.0	50.9	43.0	38.8
1931	37.1	38.4	38.4	46.5	53.5	58.4	60.6	58.2	53.8	46.6	45.5	40.6
1932	42.4	38.4	40.3	44.6	50.9	57.0	62.1	63.5	56.3	47.3	43.6	41.8
1933	36.2	39.3	44.5	48.7	54.2	60.8	65.5	64.9	60.1	50.2	42.1	35.8
1934	39.4	38.2	40.4	46.9	53.4	59.6	66.5	60.4	59.2	51.2	42.8	45.8
1935	40.0	42.6	43.5	47.1	50.0	60.5	64.6	64.0	56.8	48.6	44.2	36.4
1936	37.3	35.0	44.0	43.9	52.7	58.6	60.0	61.1	58.1	49.6	41.6	41.3
1937	40.8	41.0	38.4	47.4	54.1	58.6	61.4	61.8	56.3	50.9	41.4	37.1
1938	42.1	41.2	47.3	46.6	52.4	59.0	59.6	60.4	57.0	50.7	47.8	39.2
1939	39.4	40.9	42.4	47.8	52.4	58.0	60.7	61.8	58.2	46.7	46.6	37.8

```
> (nottem.Q <- aggregate(nottem, nfrequency = 4, FUN = mean))
```

	Qtr1	Qtr2	Qtr3	Qtr4
1920	41.933	53.100	56.133	44.400
1921	43.033	53.267	61.067	45.567
1922	38.567	51.867	55.133	43.533
1923	41.600	49.233	59.400	41.033
1924	38.367	52.133	58.467	45.933
1925	40.433	52.767	59.167	41.467
1926	42.000	52.100	60.667	42.700
1927	41.067	51.267	58.533	42.600


```

1928 41.567 51.533 59.367 43.500
1929 35.700 51.300 60.867 44.667
1930 39.967 52.833 59.567 44.233
1931 37.967 52.800 57.533 44.233
1932 40.367 50.833 60.633 44.233
1933 40.000 54.567 63.500 42.700
1934 39.333 53.300 62.033 46.600
1935 42.033 52.533 61.800 43.067
1936 38.767 51.733 59.733 44.167
1937 40.067 53.367 59.833 43.133
1938 43.533 52.667 59.000 45.900
1939 40.900 52.733 60.233 43.700

```

11 Irregularly spaced time series data

The time series data objects we have encountered so far are *regular* in the sense that the time passing from one observation to another is always the same. Some type of data (in many cases high frequency data such as data from financial markets) is irregular so that the time between observations is not always the same.

There are several packages that attempt to handle this problem, but here we will only have a brief look at one of them. This is package **zoo** (Z[eilies]’s Ordered Observations, after the original author’s last name). Instead of a **ts** object we now create a **zoo** object. As before we need a vector or matrix with data. However, we cannot just set the start or end time of the date. We need to have an index along which the data are ordered.

We start by constructing such a **zoo** object. First we simulate some observations:

```

> set.seed(1069)
> (z1.data <- rnorm(10))

[1] -1.191128 -0.092736 -0.064871 -0.717833 -0.953335 -1.360759
[7]  0.814218  0.640844 -0.525930 -3.256695

> (z2.data <- rnorm(10))

[1]  2.35276 -0.74599  0.56855  0.72583  0.67866  1.49532
[7] -1.14575  2.12064  1.53274 -0.62578

```

Then we create an index along which these observations are ordered:

```

> (z1.index <- Sys.Date() - sample(1:20, size = 10))

```

```

[1] "2011-09-03" "2011-09-05" "2011-09-07" "2011-09-04"
[5] "2011-09-18" "2011-09-19" "2011-09-21" "2011-09-12"
[9] "2011-09-14" "2011-09-17"

> (z2.index <- Sys.Date() - sample(1:20, size = 10))

[1] "2011-09-15" "2011-09-09" "2011-09-02" "2011-09-03"
[5] "2011-09-10" "2011-09-08" "2011-09-04" "2011-09-05"
[9] "2011-09-17" "2011-09-11"

```

With the data and the index we can create the `zoo` object using the function `zoo(x, order.by)`, where `x` is the data and `order.by` is the index:

```

> library(zoo)
> (z1 <- zoo(z1.data, z1.index))

2011-09-03 2011-09-04 2011-09-05 2011-09-07 2011-09-12
-1.191128 -0.717833 -0.092736 -0.064871 0.640844
2011-09-14 2011-09-17 2011-09-18 2011-09-19 2011-09-21
-0.525930 -3.256695 -0.953335 -1.360759 0.814218

> (z2 <- zoo(z2.data, z2.index))

2011-09-02 2011-09-03 2011-09-04 2011-09-05 2011-09-08
0.56855 0.72583 -1.14575 2.12064 1.49532
2011-09-09 2011-09-10 2011-09-11 2011-09-15 2011-09-17
-0.74599 0.67866 -0.62578 2.35276 1.53274

```

In real applications, of course, both data and the index are most probably available to the researcher.

There are several standard functions that become equipped with methods for `zoo` objects, such as `plot` etc.

However, combining different `zoo` objects we cannot use `ts.intersect()`, `ts.union()` etc. Instead `merge()` has methods for `zoo` objects which can be used; see `?merge.zoo` for details. We can for instance create both the intersection and union of two `zoo` objects:

```

> merge(z1, z2)

      z1      z2
2011-09-02      NA 0.56855
2011-09-03 -1.191128 0.72583
2011-09-04 -0.717833 -1.14575
2011-09-05 -0.092736 2.12064
2011-09-07 -0.064871      NA
2011-09-08      NA 1.49532

```

```

2011-09-09      NA -0.74599
2011-09-10      NA  0.67866
2011-09-11      NA -0.62578
2011-09-12  0.640844      NA
2011-09-14 -0.525930      NA
2011-09-15      NA  2.35276
2011-09-17 -3.256695  1.53274
2011-09-18 -0.953335      NA
2011-09-19 -1.360759      NA
2011-09-21  0.814218      NA

```

```
> merge(z1, z2, all = FALSE)
```

```

           z1      z2
2011-09-03 -1.191128  0.72583
2011-09-04 -0.717833 -1.14575
2011-09-05 -0.092736  2.12064
2011-09-17 -3.256695  1.53274

```

In the case of creating the union missing values are imputed at appropriate places.

Sometimes one wants to transform an irregular time series to a regular. In the previous example we have data observed on an irregular daily basis. Let us take the series `z1` as an example. There are ten observations starting October 8, 2010 and ending October 26, 2010. If we create a regular daily series with observations for each day between October 8, 2010 and October 26, 2010 (inclusive) we will have 19 observations. In reality we have 10 so days with no observation must be filled with NA's, which may later be replaced by some imputed value.

The first step to create a regular time series is to create the time index along which observations are ordered. We use the starting and ending dates of `z1` and create an index which is increased by one (day) in each step:

```
> Z1.index <- zoo(, seq(start(z1), end(z1), by = 1))
```

The length of the index is 19 and has (by definition) the start and end dates of `z1`. The second step is to merge this index with the irregular time series object `z1`:

```
> (Z1 <- merge(z1, Z1.index))
```

```

2011-09-03 2011-09-04 2011-09-05 2011-09-06 2011-09-07
-1.191128 -0.717833 -0.092736      NA -0.064871
2011-09-08 2011-09-09 2011-09-10 2011-09-11 2011-09-12
      NA      NA      NA      NA  0.640844
2011-09-13 2011-09-14 2011-09-15 2011-09-16 2011-09-17

```

```

      NA -0.525930      NA      NA -3.256695
2011-09-18 2011-09-19 2011-09-20 2011-09-21
      -0.953335 -1.360759      NA  0.814218

```

Merging fills days for which we have no observation with NA's. Once this step is completed it may be a good idea to redefine the data series to a regular time series object using the `zooreg` class:

```

> (Z1 <- zoo(Z1, frequency = 365))

2011-09-03 2011-09-04 2011-09-05 2011-09-06 2011-09-07
      -1.191128 -0.717833 -0.092736      NA -0.064871
2011-09-08 2011-09-09 2011-09-10 2011-09-11 2011-09-12
      NA      NA      NA      NA  0.640844
2011-09-13 2011-09-14 2011-09-15 2011-09-16 2011-09-17
      NA -0.525930      NA      NA -3.256695
2011-09-18 2011-09-19 2011-09-20 2011-09-21
      -0.953335 -1.360759      NA  0.814218

```

Package **zoo** comes with several alternatives to replace NA's with some numerical value:

`na.approx()` This method replaces missing values with a linear approximation from surrounding observations using the function `approx()`. Example: Say we have three observations 1, 2, 10. We want to create two additional observations between the first and the second and between the second and the third. A linear approximation would result in the series 1, 1.5, 2, 6, 10:

```

> x <- 1:3
> y <- c(1, 2, 10)
> approx(x, y, n = 5)

$x
[1] 1.0 1.5 2.0 2.5 3.0

$y
[1] 1.0 1.5 2.0 6.0 10.0

```

Applied to `Z1` we get:

```

> na.approx(Z1)

2011-09-03 2011-09-04 2011-09-05 2011-09-06 2011-09-07
      -1.191128 -0.717833 -0.092736 -0.078804 -0.064871
2011-09-08 2011-09-09 2011-09-10 2011-09-11 2011-09-12

```

```

      0.076272  0.217415  0.358558  0.499701  0.640844
2011-09-13 2011-09-14 2011-09-15 2011-09-16 2011-09-17
      0.057457 -0.525930 -1.436185 -2.346440 -3.256695
2011-09-18 2011-09-19 2011-09-20 2011-09-21
      -0.953335 -1.360759 -0.273270  0.814218

```

`na.locf` Replaces each NA with the most recent non-NA. If the first observation is NA it is removed by default:

```
> na.locf(Z1)
```

```

2011-09-03 2011-09-04 2011-09-05 2011-09-06 2011-09-07
-1.191128 -0.717833 -0.092736 -0.092736 -0.064871
2011-09-08 2011-09-09 2011-09-10 2011-09-11 2011-09-12
-0.064871 -0.064871 -0.064871 -0.064871  0.640844
2011-09-13 2011-09-14 2011-09-15 2011-09-16 2011-09-17
 0.640844 -0.525930 -0.525930 -0.525930 -3.256695
2011-09-18 2011-09-19 2011-09-20 2011-09-21
-0.953335 -1.360759 -1.360759  0.814218

```

`na.spline` Replaces NA with values from a monotone cubic function fitted to the data:

```
> na.spline(Z1)
```

```

2011-09-03 2011-09-04 2011-09-05 2011-09-06 2011-09-07
-1.191128 -0.717833 -0.092736  0.044282 -0.064871
2011-09-08 2011-09-09 2011-09-10 2011-09-11 2011-09-12
-0.031040  0.149314  0.383324  0.578123  0.640844
2011-09-13 2011-09-14 2011-09-15 2011-09-16 2011-09-17
 0.413057 -0.525930 -2.321044 -3.751878 -3.256695
2011-09-18 2011-09-19 2011-09-20 2011-09-21
-0.953335 -1.360759 -1.843246  0.814218

```

12 Disaggregation of time series data

The procedure used to create regular time series data out of irregular time series data can also be used in the case one need to disaggregate data which are standard time series objects. That is, transform data from lower frequency to higher frequency.

Consider the data set `LakeHuron` which consists of yearly observations of the level of Lake Huron 1875–1972, which we want to disaggregate to quarterly data. To make it simple we just consider the period 1875–1884:

```
> (LakeHuron <- window(LakeHuron, end = 1884))
```

Time Series:

Start = 1875

End = 1884

Frequency = 1

[1] 580.38 581.86 580.97 580.80 579.79 580.39 580.42 580.82

[9] 581.40 581.32

We then create the time index, where we now want to have an index with dates for each quarter:

```
> LH.index <- zoo(, seq(start(LakeHuron)[1], end(LakeHuron)[1],  
+      by = 1/4))
```

```
> (LH <- merge(as.zoo(LakeHuron), LH.index))
```

1875(1)	1875(2)	1875(3)	1875(4)	1876(1)	1876(2)	1876(3)	1876(4)
580.38	NA	NA	NA	581.86	NA	NA	NA
1877(1)	1877(2)	1877(3)	1877(4)	1878(1)	1878(2)	1878(3)	1878(4)
580.97	NA	NA	NA	580.80	NA	NA	NA
1879(1)	1879(2)	1879(3)	1879(4)	1880(1)	1880(2)	1880(3)	1880(4)
579.79	NA	NA	NA	580.39	NA	NA	NA
1881(1)	1881(2)	1881(3)	1881(4)	1882(1)	1882(2)	1882(3)	1882(4)
580.42	NA	NA	NA	580.82	NA	NA	NA
1883(1)	1883(2)	1883(3)	1883(4)	1884(1)			
581.40	NA	NA	NA	581.32			

The above methods can then be used to replace missing values:

```
> na.approx(LH)
```

1875(1)	1875(2)	1875(3)	1875(4)	1876(1)	1876(2)	1876(3)	1876(4)
580.38	580.75	581.12	581.49	581.86	581.64	581.41	581.19
1877(1)	1877(2)	1877(3)	1877(4)	1878(1)	1878(2)	1878(3)	1878(4)
580.97	580.93	580.88	580.84	580.80	580.55	580.29	580.04
1879(1)	1879(2)	1879(3)	1879(4)	1880(1)	1880(2)	1880(3)	1880(4)
579.79	579.94	580.09	580.24	580.39	580.40	580.40	580.41
1881(1)	1881(2)	1881(3)	1881(4)	1882(1)	1882(2)	1882(3)	1882(4)
580.42	580.52	580.62	580.72	580.82	580.97	581.11	581.25
1883(1)	1883(2)	1883(3)	1883(4)	1884(1)			
581.40	581.38	581.36	581.34	581.32			

```
> na.locf(LH)
```

1875(1)	1875(2)	1875(3)	1875(4)	1876(1)	1876(2)	1876(3)	1876(4)
580.38	580.38	580.38	580.38	581.86	581.86	581.86	581.86
1877(1)	1877(2)	1877(3)	1877(4)	1878(1)	1878(2)	1878(3)	1878(4)

```

580.97 580.97 580.97 580.97 580.80 580.80 580.80 580.80
1879(1) 1879(2) 1879(3) 1879(4) 1880(1) 1880(2) 1880(3) 1880(4)
579.79 579.79 579.79 579.79 580.39 580.39 580.39 580.39
1881(1) 1881(2) 1881(3) 1881(4) 1882(1) 1882(2) 1882(3) 1882(4)
580.42 580.42 580.42 580.42 580.82 580.82 580.82 580.82
1883(1) 1883(2) 1883(3) 1883(4) 1884(1)
581.40 581.40 581.40 581.40 581.32

```

```
> na.spline(LH)
```

```

1875(1) 1875(2) 1875(3) 1875(4) 1876(1) 1876(2) 1876(3) 1876(4)
580.38 581.17 581.65 581.87 581.86 581.69 581.43 581.16
1877(1) 1877(2) 1877(3) 1877(4) 1878(1) 1878(2) 1878(3) 1878(4)
580.97 580.90 580.91 580.90 580.80 580.56 580.24 579.95
1879(1) 1879(2) 1879(3) 1879(4) 1880(1) 1880(2) 1880(3) 1880(4)
579.79 579.82 579.99 580.21 580.39 580.47 580.47 580.44
1881(1) 1881(2) 1881(3) 1881(4) 1882(1) 1882(2) 1882(3) 1882(4)
580.42 580.46 580.55 580.67 580.82 580.98 581.14 581.28
1883(1) 1883(2) 1883(3) 1883(4) 1884(1)
581.40 581.47 581.50 581.45 581.32

```

In this case data were in levels, in which case we know start and end points. However, special care has to be taken if data measures flows where we only know the aggregate change. Such methods are not treated here.